

LLR-40: slowdown vs fastest representation per kernel
(preset M, float64, 1 thread, GNU 13.3.1; '-' = no ok timing)

